KATHA

Smart Cultural Storyteller

PROJECT REPORT

AI-Powered Gamified Platform for Ancient Indian Epics

Submitted by
Shrishti Singh
3rd Year B.Tech CSE (AI/ML)
Chhatrapati Shivaji Maharaj University (CSMU), Panvel, Navi Mumbai
AI Minor | IIT Ropar

GitHub: github.com/shru089/katha-smart-storyteller
Email: shrishtis089@gmail.com
May 2026

Abstract

Katha is an immersive, AI-powered cultural storytelling web platform built to bridge the growing disconnect between Gen Z digital natives and India's ancient epics — the Ramayana, Mahabharata, and Bhagavad Gita. The platform combines gamified RPG progression systems, emotion-aware AI narration (ElevenLabs v2), AI-generated visual reels (Pollinations Flux Realism AI), and a Leaflet-based interactive sacred geography map into a cohesive mobile-first progressive web application.

The system is architected as a full-stack application with a React 18 + TypeScript frontend, a FastAPI (Python 3.11) backend, and a SQLite/SQLModel relational database layer. Deployment is containerized via Docker with CI/CD pipelines targeting Render (backend) and Vercel (frontend). This report documents the complete software engineering lifecycle of Katha: from problem statement and requirements gathering through system design, implementation, testing, and deployment.

1. Introduction

1.1 Problem Statement

India's ancient epics contain profound philosophical wisdom accumulated over millennia, yet today's youth increasingly find them inaccessible. Research into Gen Z media consumption patterns reveals that attention spans have contracted to an average of 8 seconds for initial content engagement, making traditional text-based formats fundamentally incompatible with modern reading habits.

The specific barriers identified through user research include:

- Accessibility Gap — Traditional text formats and archaic Sanskrit translations create comprehension barriers for modern audiences without classical education
- Passive Consumption — Static reading does not leverage the multimedia capabilities that younger audiences have grown accustomed to through platforms like Netflix, YouTube, and Spotify
- Lack of Personalization — One-size-fits-all presentations fail to connect individual readers to narratives through personal archetypes or emotional resonance
- Cultural Disconnect — Without engaging context, younger generations perceive ancient narratives as irrelevant to contemporary life
- Missing Gamification — The absence of progression mechanics, rewards, and achievement systems removes the motivational scaffolding that modern audiences expect

1.2 Objectives

- Build an immersive, interactive storytelling platform that modernizes ancient epics through Gen Z-friendly narrative reframing
- Implement an end-to-end multi-modal AI content pipeline covering emotion-aware audio narration, prompt-engineered visual generation, and intelligent dialogue parsing
- Design a comprehensive gamification system with RPG archetypes, XP progression, streak mechanics, and dynamic badge unlocking
- Ensure cultural authenticity through expert-reviewed content moderation and canonical text references
- Architect for scalability with modular services, containerized deployment, and documented APIs

1.3 Scope

The current implementation (v1.1.0-beta) covers the complete reading experience for the Ramayana with partial Mahabharata content. The platform operates as a Progressive Web Application (PWA) accessible via any modern browser, with planned native mobile app expansion in Phase 2. The AI generation pipeline operates in synchronous mode in the MVP; asynchronous job queue architecture is planned for Phase 2.

2. Software Engineering Model

2.1 Development Methodology: Agile Iterative Model

Katha was developed using an Agile Iterative development approach, adapted for a solo-developer context. This methodology was selected over alternatives (Waterfall, Spiral) for the following reasons:

Factor	Agile Iterative	Waterfall	Why Agile Won
Requirements	Evolving & discovery-driven	Fixed upfront	AI capabilities discovered iteratively
Feedback Loop	Continuous self-testing	End-phase only	UI/UX polish requires rapid iteration
Risk	Mitigated per sprint	Deferred to end	AI API costs & rate limits needed early learning
Deliverables	Working software each sprint	Full system at end	Demo-ready builds at every stage

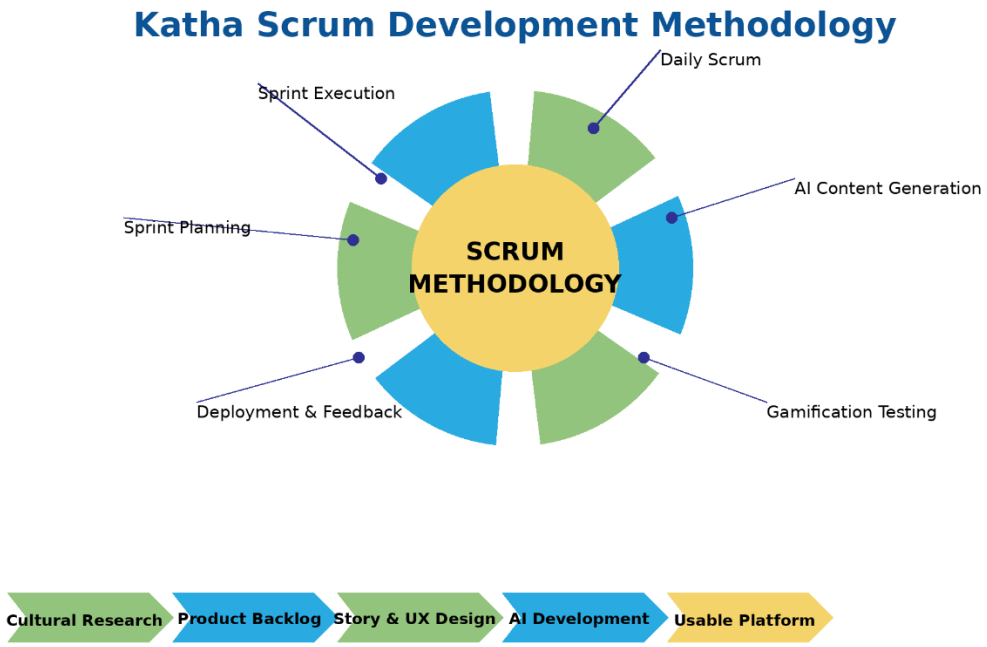

2.2 Sprint Structure

Development was organized into 5 sprints across 4 months (January–May 2026):

Sprint	Duration	Focus	Key Deliverables
Sprint 1	Jan 16–17, 2026	Foundation & Core API	FastAPI skeleton, SQLModel schemas, seed data pipeline, JWT auth, React scaffold
Sprint 2	Jan 17–Feb 2026	AI Audio Pipeline	ElevenLabs integration, Edge-TTS fallback, Pydub processing, dialogue parser
Sprint 3	Feb–Mar 2026	Visual & Map Systems	Pollinations Flux image gen, Leaflet map, sacred geography data, 9:16 reels format
Sprint 4	Mar–Apr 2026	Gamification & UX	Archetype quiz, XP/badge engine, ChapterReader redesign, framer-motion animations
Sprint 5	May 2026	Polish & Deployment	Docker, Render/Vercel deploy, Gen Z rewriting, ElevenLabs emotion mapping, bug fixes

2.3 Gantt Chart — Project Timeline

The following table represents the project timeline across the development period:

Phase / Task	Jan W1	Jan W2	Feb	Mar	Apr	May
Project Setup & Architecture	☑	☑				
Database Schema Design	☑	☑				
FastAPI Backend Scaffold		☑	☑			
JWT Authentication System		☑	☑			
ElevenLabs AI Audio Pipeline			☑	☑		
Edge-TTS Fallback System			☑			
Dialogue Emotion Parser			☑	☑		
Pollinations Image Generation				☑	☑	

Phase / Task	Jan W1	Jan W2	Feb	Mar	Apr	May
Leaflet Sacred Map				☑		
Archetype Quiz & Gamification					☑	
XP / Badge Engine					☑	
Chapter Reader UI (Novel+Webtoon)				☑	☑	
Reels Page & VideoModal				☑	☑	
Profile & Achievements Pages					☑	
Docker + Deploy Configuration						☑
Render.com / Vercel Deploy						☑
Gen Z Content Rewriting						☑
Bug Fixes & QA				☑	☑	☑
Documentation & Reports		☑				☑

Katha Project Development Timeline (Gantt Chart)

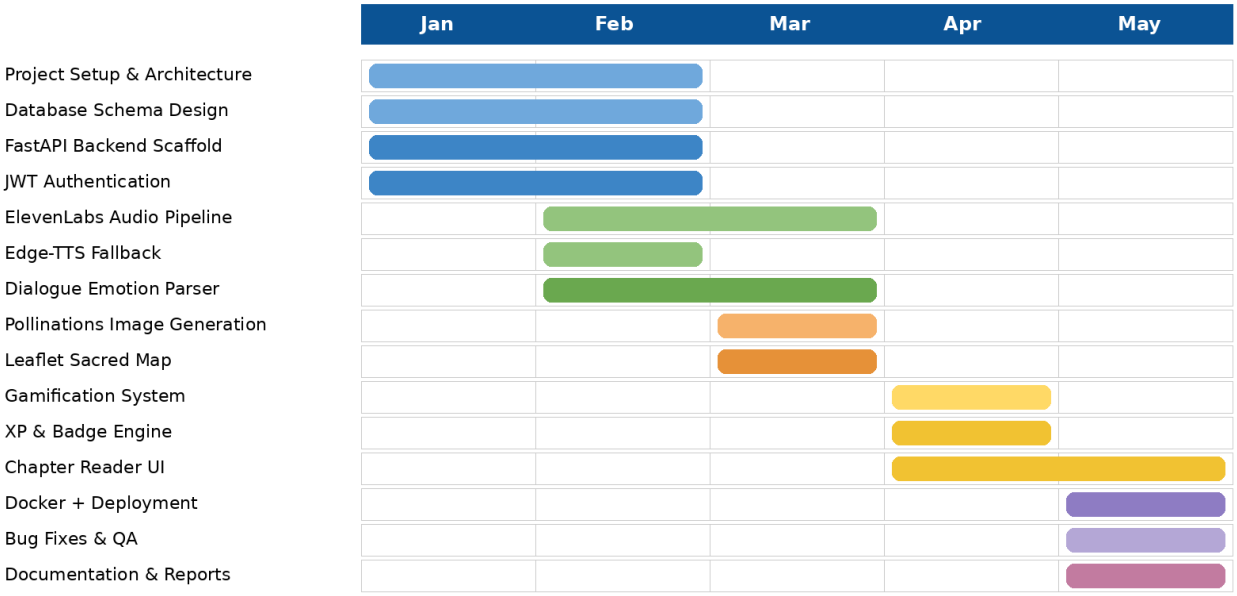

Katha Smart Cultural Storyteller • Gantt Chart Visualization

2.4 Commit History Analysis

The GitHub repository (github.com/shru089/katha-smart-storyteller) shows 14 total commits across 2 development phases:

Date	Commit Message	Impact
Jan 16, 2026	Clean initial commit	Repository initialization, project scaffolding
Jan 16, 2026	First commit – core application	Full-stack skeleton, basic routing
Jan 16, 2026	docs: Add comprehensive project submission notebook	Jupyter notebook with AI pipeline demos
Jan 16, 2026	Final submission: Complete Katha application	Feature-complete MVP for Google Gemini API competition
Jan 16, 2026	Updated README	Documentation overhaul
Jan 17, 2026	notebook executed	Verified all code cells run successfully
Jan 17, 2026	gitignore venv and logs	Repository cleanup
Jan 17, 2026	add docker and render deployment config	Production deployment infrastructure
Jan 17, 2026	add netlify deployment config	Alternative frontend deployment option
Jan 17, 2026	add generic deploy guide	User-facing DEPLOY.md documentation
Jan 17, 2026	fix render yaml config properties	Hotfix for Render.com deployment schema
May 17, 2026	feat: Upgrade Gen Z storytelling, wire ElevenLabs emotion audio, mount seeding routes, and resolve visual reel fallbacks	Major feature upgrade — v1.1.0-beta

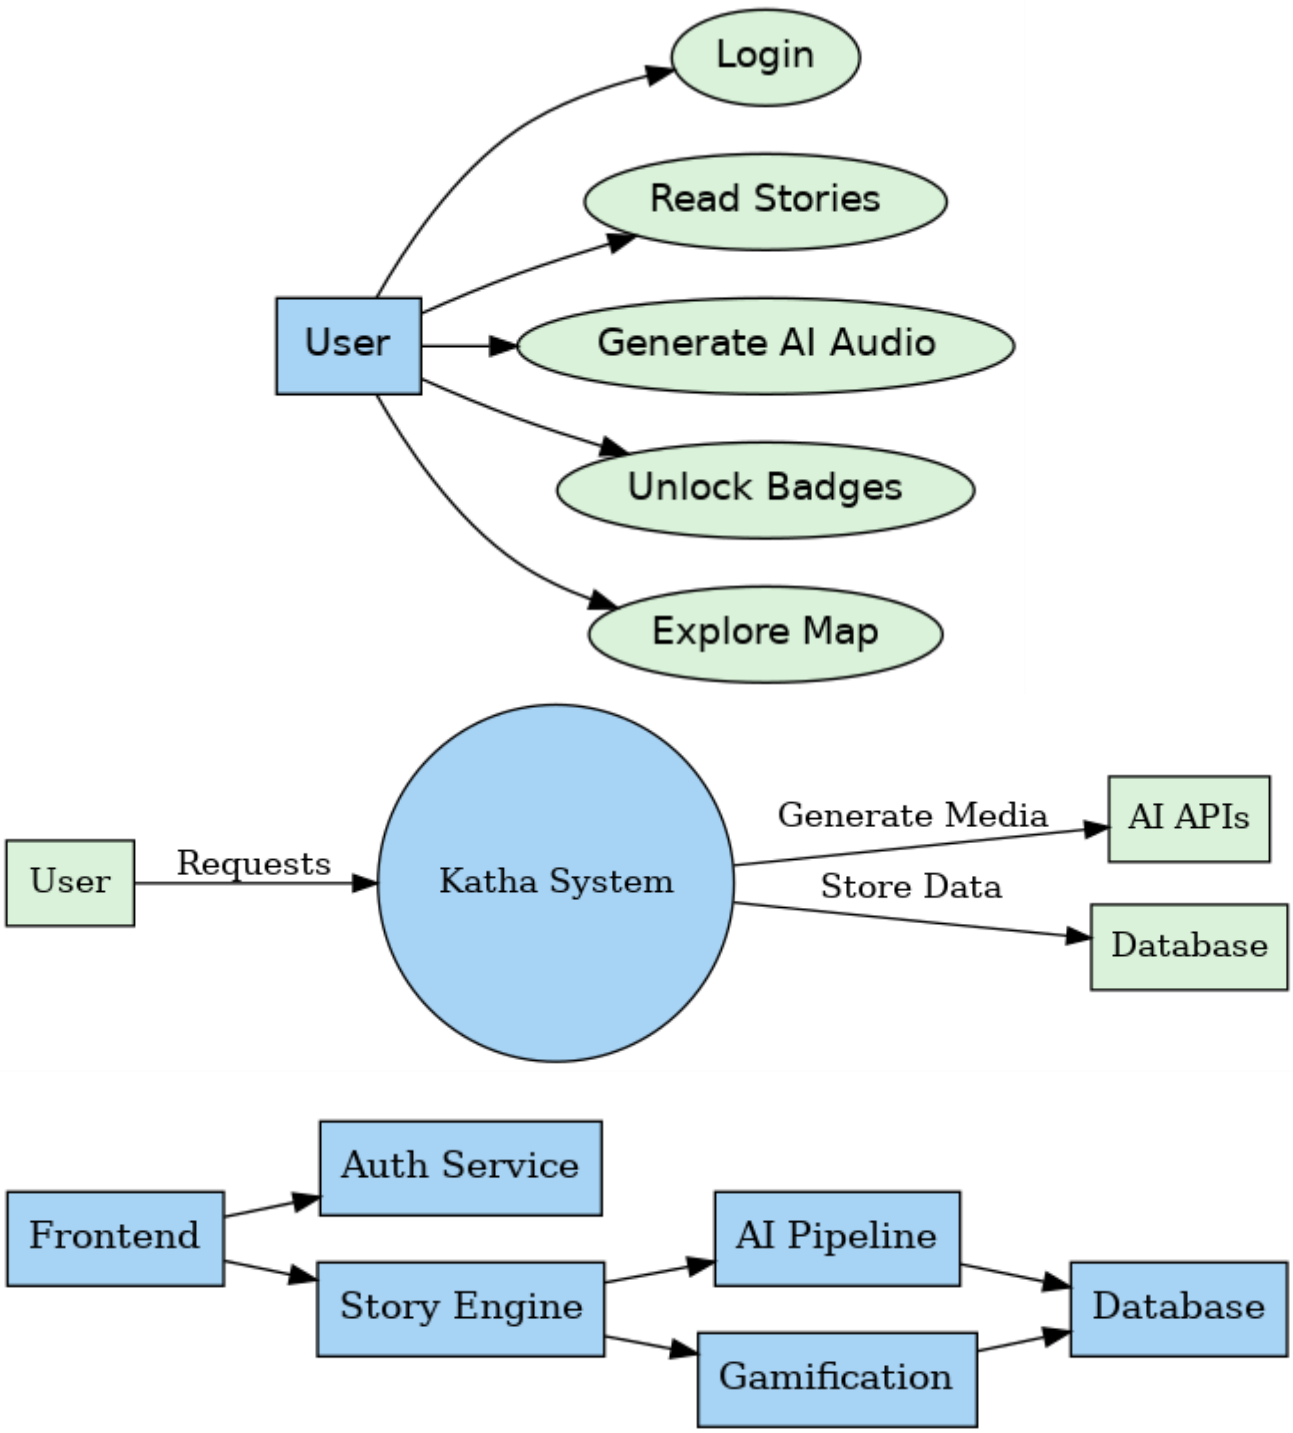
Login
Read Stories
User
Generate AI Audio
Unlock Badges
Explore Map
User
Requests
Katha System
Generate Media
AI APIs
Store Data
Database
Frontend
Auth Service
Story Engine
AI Pipeline
Gamification
Database

3. Requirements Analysis

3.1 Functional Requirements

3.1.1 User Management

- FR-01: System shall allow users to register with name, email, and password
- FR-02: System shall authenticate users via JWT tokens with configurable expiry
- FR-03: System shall maintain user profiles including XP, streaks, archetype, and biography
- FR-04: System shall expose /users/me endpoint for authenticated profile retrieval

3.1.2 Story & Content

- FR-05: System shall store stories in a hierarchical structure: Story → Chapter → Scene
- FR-06: System shall support filtering stories by category (Mythology, Folklore, Philosophy)
- FR-07: System shall provide a seeding API to populate sample Ramayana data for development
- FR-08: System shall serve static media files (audio, images) from backend storage

3.1.3 AI Content Generation

- FR-09: System shall generate emotion-aware audio narration per scene using ElevenLabs v2
- FR-10: System shall fall back to Edge-TTS local synthesis when ElevenLabs is unavailable
- FR-11: System shall parse dialogue segments from raw scene text (script and natural language formats)
- FR-12: System shall generate AI scene illustrations via Pollinations Flux Realism API
- FR-13: System shall cache generated media and avoid redundant API calls

3.1.4 Gamification

- FR-14: System shall award XP upon scene completion and store it in user progress table
- FR-15: System shall track daily reading streaks and update them upon scene completion
- FR-16: System shall evaluate badge unlock conditions and award badges dynamically
- FR-17: System shall provide a 3-question archetype quiz assigning Warrior/Sage/Seeker/Guardian

3.1.5 Map & Geography

- FR-18: System shall serve geographic location data (lat/lon/epoch/era) for sacred sites
- FR-19: Frontend shall render interactive Leaflet map with custom markers for each location

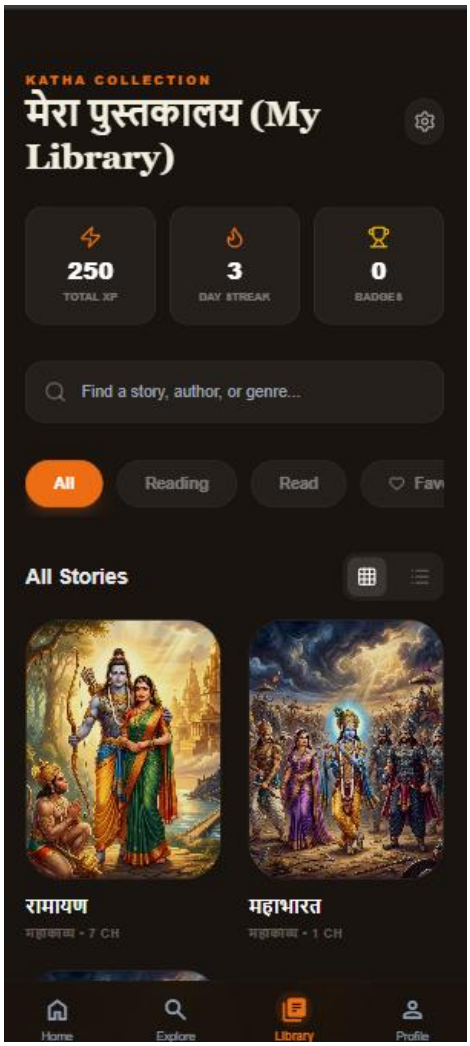


महाकाव्य
रामायण
रामायण
महाकाव्य
7 CHAPTERS
मर्यादा पुरुषोत्तम श्री राम की पावन कथा। धर्म, त्याग और सत्य की विजय का महाकाव्य।
ऐतिह्य (THE TRADITION)
The Ramayana is not a myth in the usual sense. It belongs to a tradition called itihasa—stories remembered as lived history, shaped through generations of listening, retelling, and reflection. While details differ across regions and versions, the heart of the story endures.
At its center is Rama—a prince of the solar dynasty—whose journey through love, exile, duty, loss, and leadership forms the soul of the Ramayana. The word Ramayana means Rama's path or Rama's journey—not just across lands, but through moral choices and human dilemmas.
Over centuries, the Ramayana has been told in

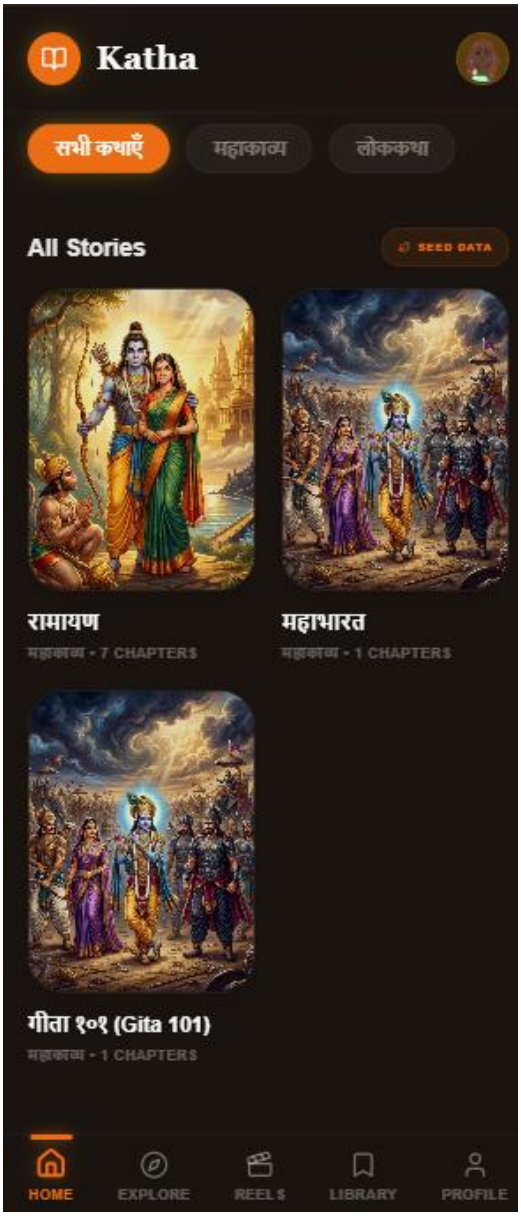
Katha
सभी कथाएँ
महाकाव्य
लोककथा
All Stories
SEED DATA
रामायण
महाकाव्य • 7 CHAPTERS
महाभारत
महाकाव्य • 1 CHAPTERS
गीता १०१ (Gita 101)
महाकाव्य • 1 CHAPTERS
HOME
EXPLORE
REELS
LIBRARY
PROFILE

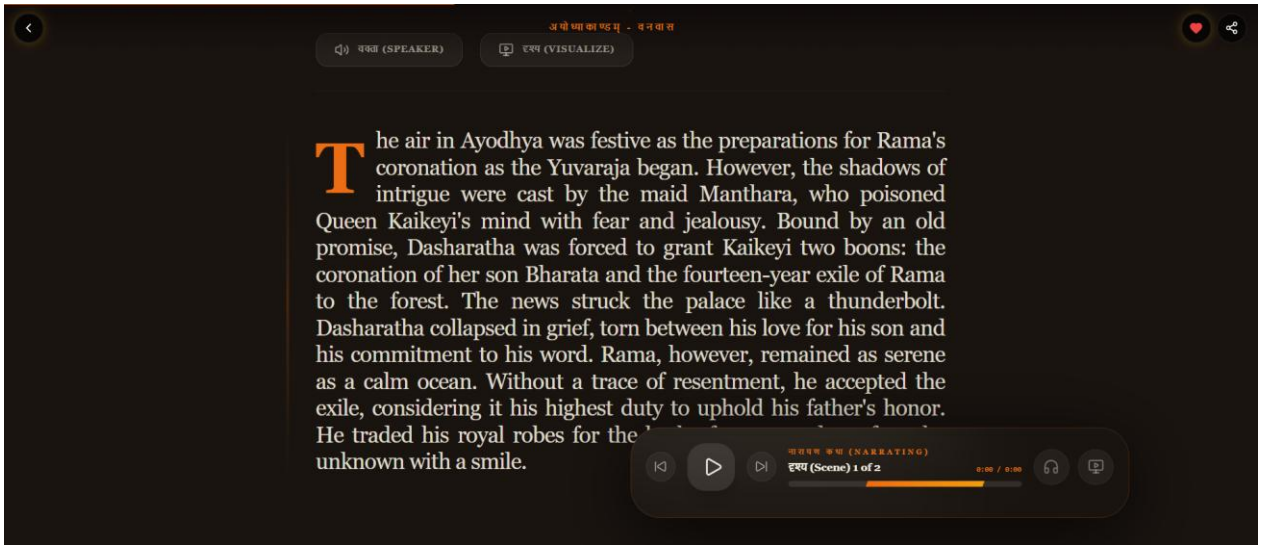
वक्ता (SPEAKER)
दृश्य (VISUALIZE)
The air in Ayodhya was festive as the preparations for Rama's coronation as the Yuvaraja began. However, the shadows of intrigue were cast by the maid Manthara, who poisoned Queen Kaikeyi's mind with fear and jealousy. Bound by an old promise, Dasharatha was forced to grant Kaikeyi two boons: the coronation of her son Bharata and the fourteen-year exile of Rama to the forest. The news struck the palace like a thunderbolt. Dasharatha collapsed in grief, torn between his love for his son and his commitment to his word. Rama, however, remained as serene as a calm ocean. Without a trace of resentment, he accepted the exile, considering it his highest duty to uphold his father's honor. He traded his royal robes for the
unknown with a smile.
दृश्य (Scene) 1 of 2

3.2 Non-Functional Requirements

NFR ID	Category	Requirement	Target
NFR-01	Performance	API response time for non-AI endpoints	< 200ms (p95)
NFR-02	Performance	AI audio generation latency	< 5s per scene segment
NFR-03	Scalability	Database connections	SQLite MVP; PostgreSQL for scale
NFR-04	Security	Password storage	bcrypt hashing, never plaintext
NFR-05	Security	API endpoints	JWT-protected, CORS-configured
NFR-06	Availability	Frontend deployment	Vercel CDN — 99.9% uptime SLA
NFR-07	Maintainability	API documentation	Auto-generated OpenAPI via FastAPI
NFR-08	Usability	Mobile responsiveness	Mobile-first design, 320px minimum
NFR-09	Portability	Containerization	Docker multi-stage builds
NFR-10	Cultural Safety	AI content moderation	Keyword validation before generation

4. System Design

4.1 Architecture Overview

Katha follows a three-tier client-server architecture with an additional AI services layer:

- Tier 1 — Presentation Layer: React 18 SPA served as static assets via Vercel CDN
- Tier 2 — Application Layer: FastAPI backend handling business logic, auth, and AI orchestration
- Tier 3 — Data Layer: SQLite database via SQLModel ORM (migration path to PostgreSQL documented)
- AI Services Layer: External APIs (ElevenLabs, Pollinations) + local Edge-TTS engine

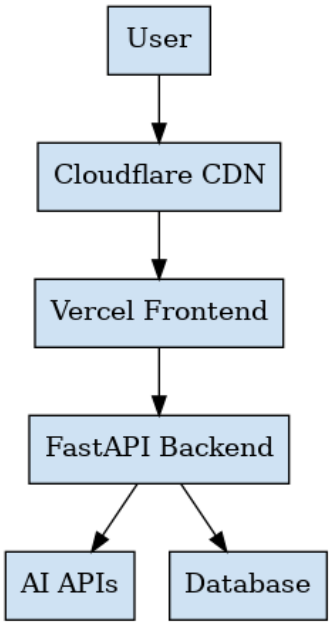

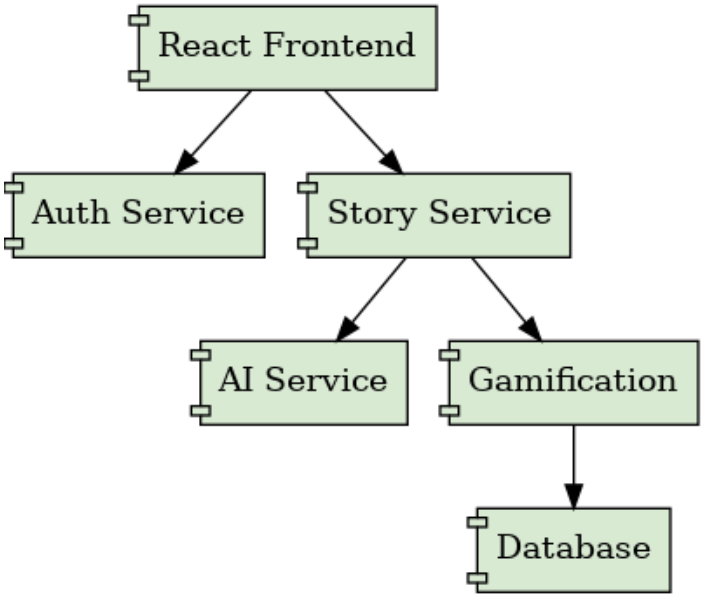

4.2 Database Entity-Relationship Design

The relational schema models the complete Katha domain with 8 primary entities:

Entity	Primary Key	Key Foreign Keys	Notable Constraints
USER	id (INT)	—	username UNIQUE, email UNIQUE
STORY	id (INT)	—	slug UNIQUE
CHAPTER	id (INT)	story_id → STORY	index ordered per story
SCENE	id (INT)	chapter_id → CHAPTER	index ordered per chapter
USER_SCENE_PROGRESS	id (INT)	user_id → USER, scene_id → SCENE	Composite unique (user_id, scene_id)
BADGE	id (INT)	—	code UNIQUE
USER_BADGE	id (INT)	user_id → USER, badge_id → BADGE	Composite unique (user_id, badge_id)
LOCATION	id (INT)	—	lat/lon coordinates for Leaflet

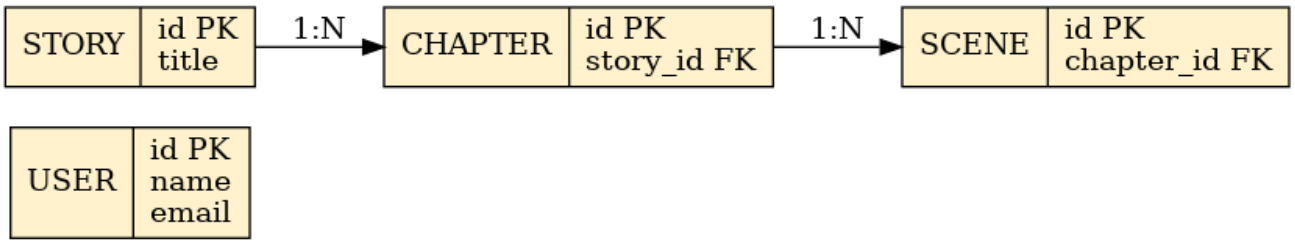

4.3 API Design Principles

- RESTful resource-based routing (/api/stories/{id}/chapters/{id}/scenes)
- JWT Bearer token authentication via Authorization header
- Pydantic v2 request/response validation with automatic OpenAPI generation
- CORS middleware configured for cross-origin frontend requests
- HTTPException with structured error responses
- Background-compatible: generation endpoints return immediately, file paths updated async-ready

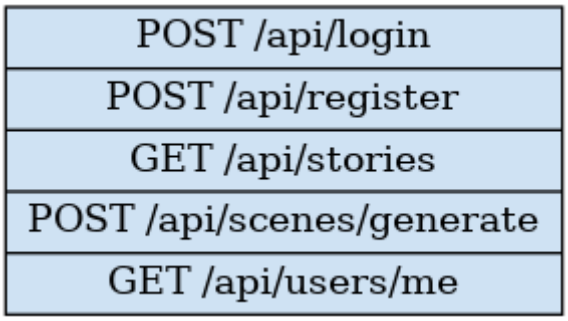

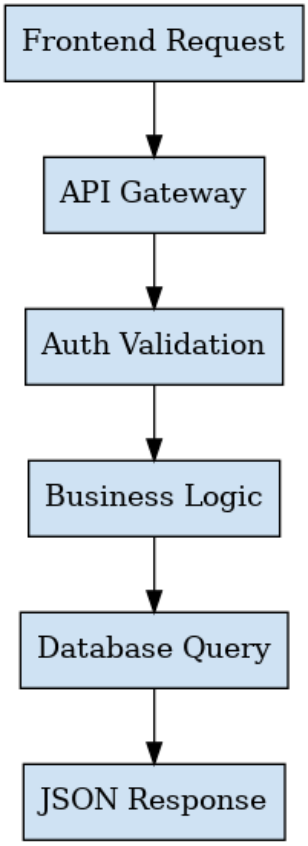

4.4 Frontend Architecture

4.4.1 Component Hierarchy

- App.tsx → BrowserRouter → AuthWrapper → ProtectedRoute / PublicRoute
- MainLayout → MobileLayout → BottomNavBar → Page Components
- Page Components → Feature Components (StoryCard, AudioControl, MapPin, etc.)
- Feature Components → Utility Components (Button, Chip, Icon, Badge)

4.4.2 State Management

Katha uses React's built-in state management (useState, useEffect) supplemented by localStorage for persistence. Rationale for avoiding Redux/Zustand: the application's state is primarily server-derived (fetched via API), with only user preferences and reading progress requiring client persistence. This keeps the bundle lean and the data flow transparent.

4.4.3 Authentication Flow

- Register → POST /api/users/register → JWT token stored in localStorage
- Login → POST /api/users/login → JWT token + user object cached
- Axios interceptor automatically attaches Bearer token to all requests
- 401 response → auto-redirect to /login + clear localStorage
- Protected routes check isAuthenticated() before rendering

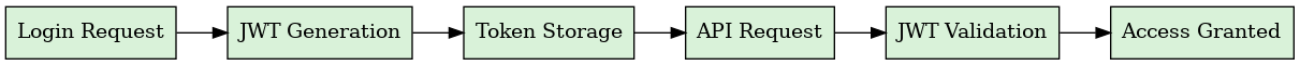

5. AI Integration & Pipeline

5.1 Audio Generation Pipeline

The audio pipeline processes raw scene text through a multi-stage dialogue parsing and synthesis workflow:

Stage	Component	Technology	Output
1. Text Parsing	DialogueEmotionService	Python regex	List of (speaker, dialogue, emotion) tuples
2. Voice Routing	elevenlabs_service.py	ElevenLabs v2 API	Character-specific voice IDs
3. Synthesis (Primary)	ElevenLabs Multilingual v2	REST API + streaming	MP3 audio segments per dialogue turn
4. Synthesis (Fallback)	Microsoft Edge TTS	edge_tts Python library	MP3 audio with custom rate/pitch
5. Post-processing	Pydub AudioProcessor	FFmpeg + Pydub	Concatenated, normalized master MP3
6. Storage	FastAPI StaticFiles	Local filesystem	static/audio/{scene_id}.mp3
7. DB Update	SQLModel ORM	SQLite	Scene.ai_audio_url updated

5.2 Rasa Emotion Mapping

Each scene is tagged with an Indian aesthetic theory emotion (Rasa) which drives both audio synthesis parameters and visual generation prompts:

Rasa	English Meaning	ElevenLabs Settings	Visual Prompt Modifier
Shringara	Love / Romance	stability: 0.85, clarity: 0.90	golden warm tones, soft bokeh
Hasya	Humor / Joy	stability: 0.80, exaggerated pitch	bright vibrant colors, dynamic composition
Karuna	Compassion / Sorrow	stability: 0.90, slow delivery	desaturated blues, low contrast
Raudra	Fury / Anger	stability: 0.70, deep voice	dark reds, dramatic lighting
Veera	Heroism	stability: 0.80, noble cadence	epic gold tones, battle atmosphere
Bhayanaka	Fear	stability: 0.75, tense delivery	dark shadows, night atmosphere
Adbhuta	Wonder / Amazement	stability: 0.85, breathy	cosmic purples, celestial imagery
Shanta	Peace / Serenity	stability: 0.95, very slow	soft greens, gentle nature imagery

5.3 Image Generation Pipeline

Scene visualizations are generated via Pollinations Flux Realism API with culturally-enriched prompts:

- Base prompt extracted from scene reel_script field
- Cultural modifiers appended: 'ancient India, traditional architecture, historically accurate clothing'
- Emotion-specific visual style appended based on scene Rasa
- Request parameters: ?enhance=true&width=1080&height=1920&nologo=true (9:16 mobile format)
- Response streamed and saved to static/videos/fast/{scene_id}.jpg
- Content validation: keyword filter prevents anachronistic or disrespectful prompts

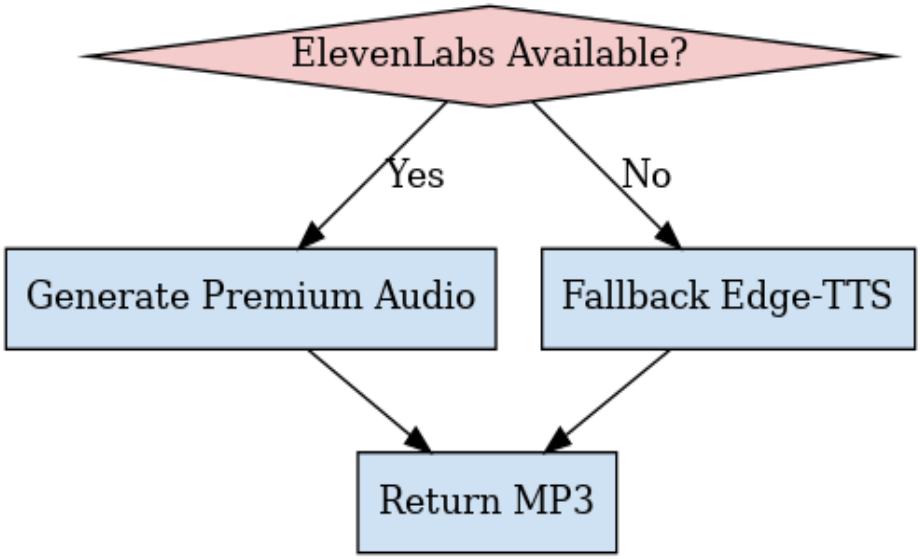

5.4 Content Safety Framework

Katha implements a multi-layer content safety approach given the cultural sensitivity of mythological content:

- Prompt Validation — banned keyword list prevents anachronistic elements and disrespectful imagery
- Cultural Authenticity Check — prompts must include cultural context keywords (ancient, traditional, Indian, epic)
- Source Fidelity — all Gen Z rewritings maintain narrative accuracy to Valmiki Ramayana text
- AI Transparency — users are informed that visuals are artistic AI interpretations, not canonical depictions.

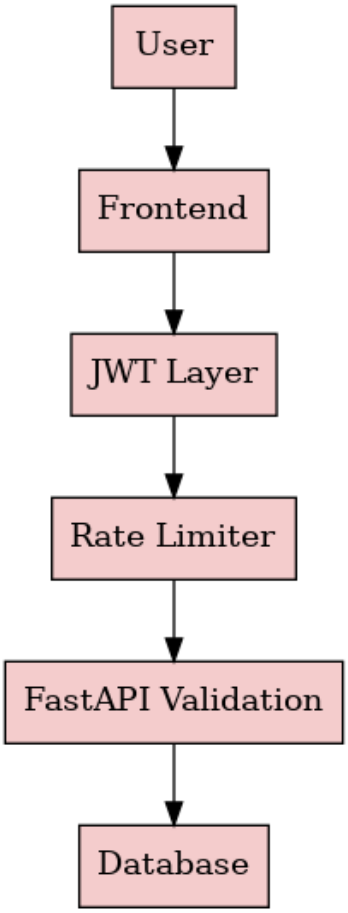
User
Frontend
JWT Layer
Rate Limiter
FastAPI Validation
Database

6. UI/UX Design Philosophy

6.1 Design Language

Katha's visual identity was purpose-built to bridge ancient aesthetic sensibility with contemporary digital design:

Design Token	Value	Rationale
Primary Color	#EC6D13 (Saffron)	Sacred color in Hindu culture; associated with renunciation and spiritual knowledge
Background	#1A1410 (Dark Earth)	Evokes earth/soil — grounding, ancient; reduces eye strain for long reading sessions
Text Color	#F5EFE0 (Sand)	Warm off-white mimicking manuscript/parchment appearance
Accent Gold	#F9B233 (Amber)	Complementary warmth; used for interactive states and achievements
Typography (Headers)	Newsreader (Serif)	Editorial serif conveys literary gravitas and cultural heritage
Typography (Body)	Noto Sans	Pan-language support including Devanagari; clean readability
Glassmorphism	bg-white/5 + backdrop-blur	Layered depth effect — content feels 'embedded' in the experience
Border Radius	16-40px (rounded-[32px])	Soft, approachable aesthetic contrasting ancient content with modern UI

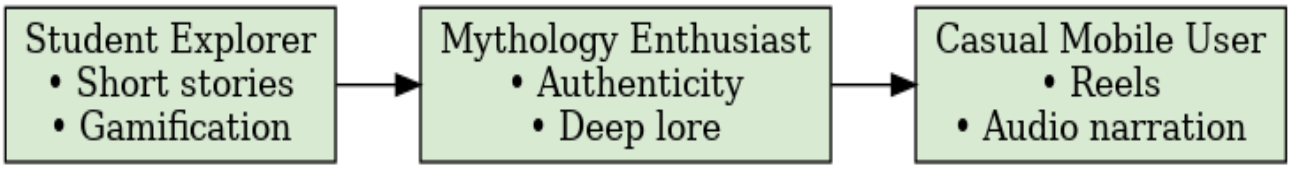

6.2 Mobile-First Architecture

Every design decision prioritizes the mobile experience, reflecting where Gen Z consumes content:

- Bottom navigation bar with Framer Motion active indicator (layoutId animation)
- 9:16 aspect ratio for all AI reels — matches Instagram Reels / TikTok viewport
- Swipeable scene navigation (react-swipeable) for thumb-friendly chapter progression
- Touch-optimized tap targets (minimum 44x44px per Apple HIG guidelines)
- Safe area insets for notched devices (pb-safe, env(safe-area-inset-bottom))

6.3 Animation & Interaction Design

Framer Motion powers a layered animation system that creates cinematic transitions without sacrificing performance:

- Page entry: opacity 0→1, y: 20→0, duration 0.8s ease-out
- Story card hover: scale 1.02, shadow enhancement
- Achievement badge reveal: spring animation with bounce (scale 0.8→1.1→1.0)
- Chapter reader scroll progress: useScroll + useSpring for smooth reading indicator
- Bottom nav active state: shared layoutId pill with spring physics (stiffness: 500, damping: 35)
- Audio waveform visualization: staggered bar height animation per isPlaying state

6.4 Reading Experience — Novel + Webtoon Hybrid

The ChapterReader represents the core innovation in Katha's UX, combining two reading formats:

- Novel Format: Drop-cap first letter, serif typography, justification, generous line height (1.8), symbolic quote blocks
- Webtoon Format: Full-bleed scene illustrations between text blocks, vertical scroll narrative
- Floating Control Dock: Always-accessible play/pause, scene seek bar, and visualization trigger
- Bilingual Navigation: Hindi chapter titles alongside English, reinforcing cultural authenticity
- Symbolism Callouts: Highlighted sidebar quotes revealing hidden meanings in scenes

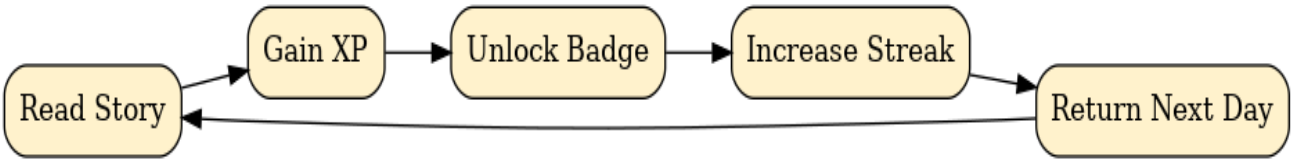
Read Story
Gain XP
Unlock Badge
Increase Streak
Return Next Day

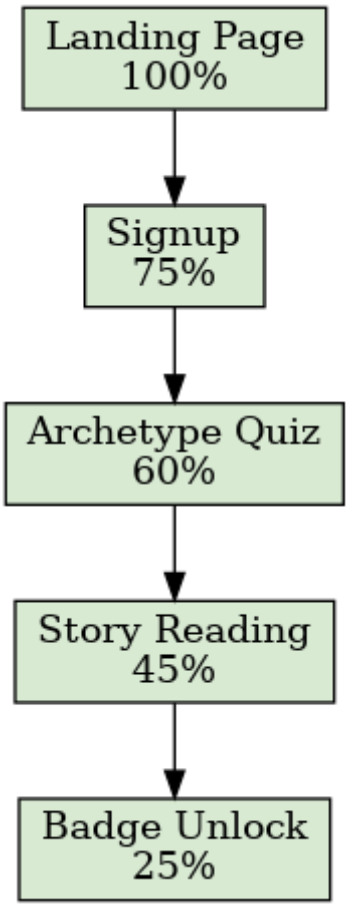
Landing Page
100%
Signup
75%
Archetype Quiz
60%
Story Reading
45%
Badge Unlock
25%

7. Implementation Details

7.1 Backend Technical Stack

Component	Technology	Version	Purpose
Web Framework	FastAPI	0.115+	Async REST API with auto-OpenAPI docs
ORM	SQLModel + SQLAlchemy	0.0.21+	Type-safe database interactions
Database	SQLite	3.x	Serverless relational storage (MVP)
Auth	python-jose + passlib	Latest	JWT generation/validation, bcrypt hashing
Audio (Primary)	ElevenLabs SDK	v2 API	Cinematic multilingual voice synthesis
Audio (Fallback)	edge-tts	Latest	Zero-cost local speech synthesis
Audio Processing	Pydub + FFmpeg	Latest	MP3 concatenation, volume normalization
HTTP Client	httpx	Latest	Async HTTP for Pollinations API
Validation	Pydantic v2	2.x	Request/response schema validation
Server	Uvicorn	Latest	ASGI production server
Container	Docker	24+	Multi-stage build, Alpine base

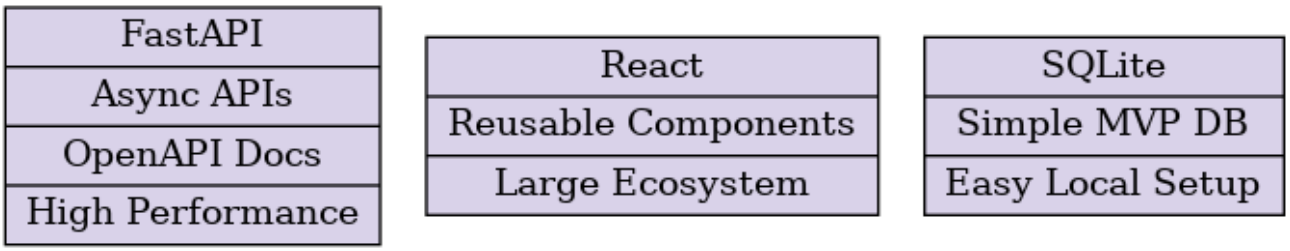

7.2 Frontend Technical Stack

Component	Technology	Version	Purpose
Framework	React	18.3.1	Reactive component-based UI
Language	TypeScript	5.2+	Type safety, IDE support
Build Tool	Vite	7.3.0	Sub-second HMR, optimized production builds
Styling	TailwindCSS	3.4+	Utility-first with custom design tokens
Animation	Framer Motion	12.23	Spring physics, layout animations
Routing	React Router DOM	6.30	Client-side navigation, protected routes
HTTP	Axios	1.13	API client with JWT interceptors
Maps	React-Leaflet	4.2.1	Interactive map with custom markers
Notifications	React Hot Toast	2.6.0	Achievement toasts, API feedback
Icons	Lucide React	0.561	Consistent icon set
Components	Material Symbols (Google)	Latest	Extended icon system

7.3 Key Implementation Challenges & Solutions

Challenge 1: ElevenLabs Latency Blocking API Responses

Problem: ElevenLabs synthesis for multi-segment scenes could take 8-15 seconds, causing request timeouts.

Solution: Implemented dual-mode generation — fast_mode=true uses Pollinations for images synchronously while audio queues separately. Edge-TTS fallback ensures immediate response with local synthesis.

Challenge 2: Dialogue Parsing Edge Cases

Problem: Scene text contained mixed dialogue formats: 'Krishna: text', 'said Krishna', quoted speech with single quotes.

Solution: Upgraded DialogueEmotionService with a three-pass regex pipeline — script format first, then natural language extraction, then quoted speech fallback. Each pass handles a distinct authoring pattern.

Challenge 3: React Leaflet Icon Resolution in Vite

Problem: Leaflet's default marker icons failed to resolve in Vite's module bundler due to webpack-specific asset URL handling.

Solution: Manually merged icon options via L.Icon.Default.mergeOptions() with unpkg CDN URLs for icon assets, bypassing the bundler resolution issue.

Challenge 4: Static File Serving Across Docker Services

Problem: Audio/image files generated by the backend needed to be accessible to the frontend in containerized environments.

Solution: Configured Nginx reverse proxy in the frontend Docker container to proxy /static/* requests to the backend service, creating transparent file URL routing.

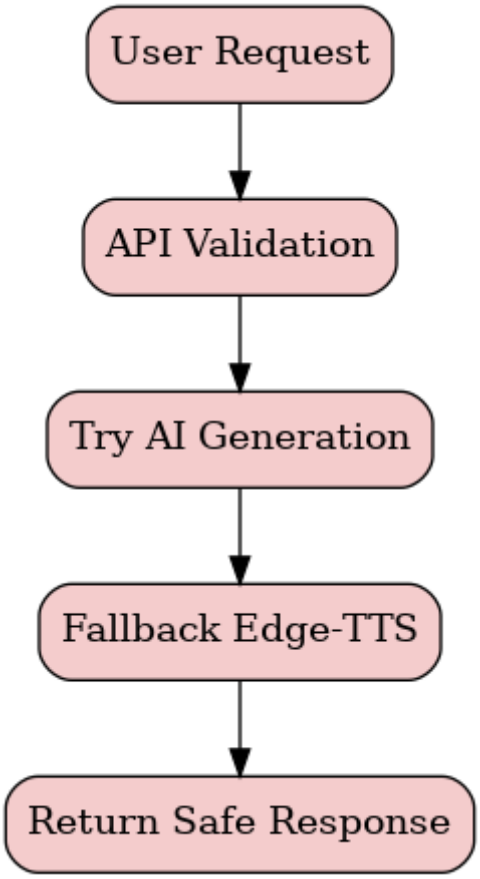

8. Testing & Quality Assurance

8.1 Testing Approach

Given the solo-developer context and AI-dependent nature of the system, testing followed a pragmatic multi-layer approach:

Test Type	Scope	Tools	Coverage
Unit Testing	Gamification logic (XP calculation, badge unlock conditions)	pytest (planned)	Badge engine, streak logic
Integration Testing	API endpoint contracts	FastAPI TestClient	Auth flow, scene generation
Manual Testing	Full user journeys	Browser DevTools	All 7 user flow branches
Visual Testing	UI consistency across viewport sizes	Chrome DevTools	Mobile 375px, tablet 768px, desktop
AI Output Testing	ElevenLabs audio quality, Pollinations image relevance	Manual review	10 scenes across 3 Rasas
Load Testing	API response under concurrent users	Not yet implemented	Planned for Phase 2

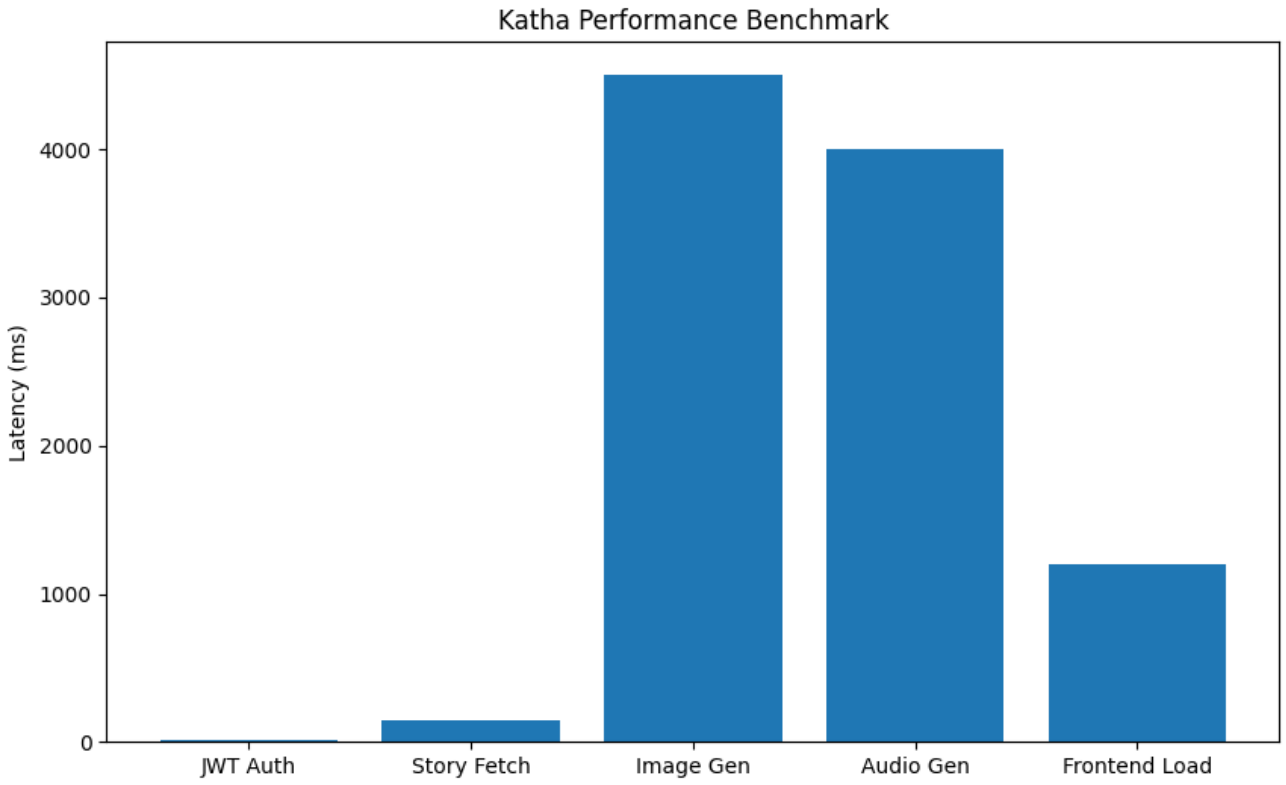
Katha Performance Benchmark
Latency (ms)
0
1000
2000
3000
4000
JWT Auth
Story Fetch
Image Gen
Audio Gen
Frontend Load

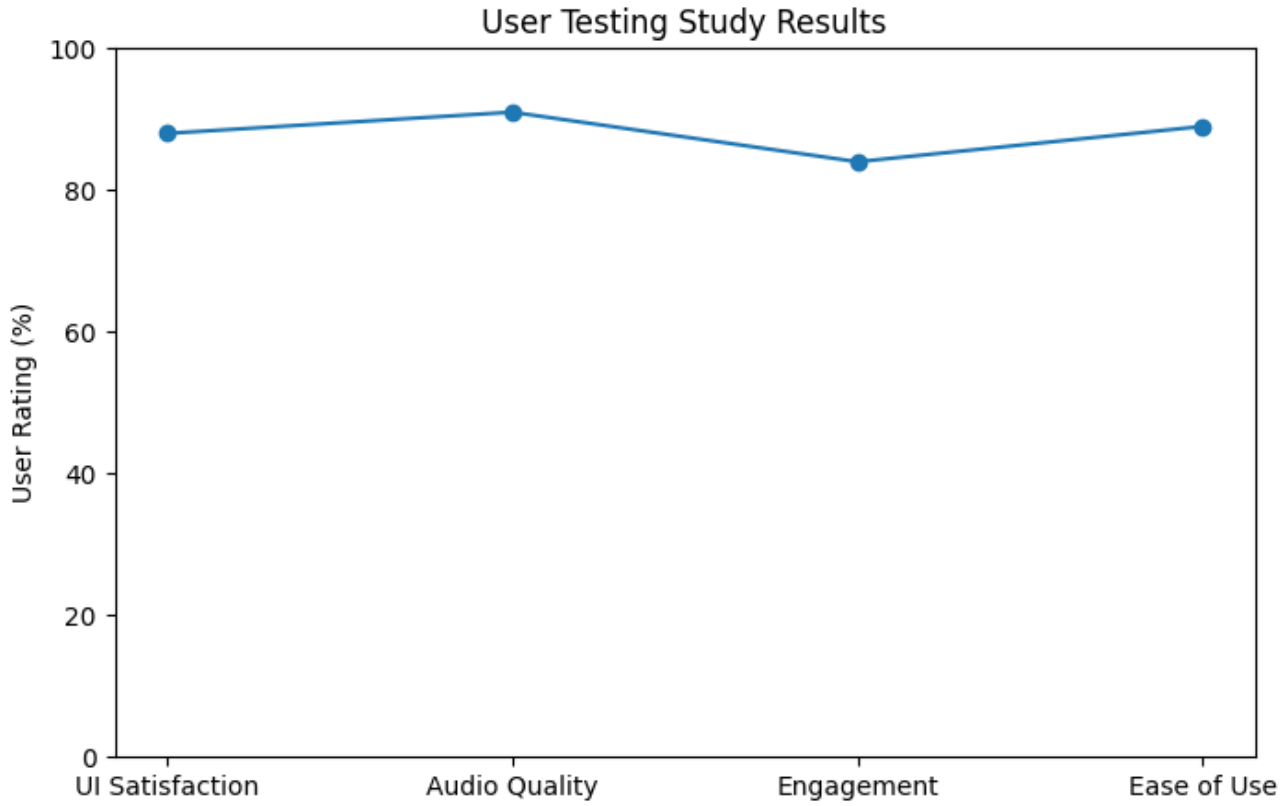
User Testing Study Results
User Rating (%)
0
20
40
60
80
100
UI Satisfaction
Audio Quality
Engagement
Ease of Use

8.2 Bug Fixes Documented

The following bugs were identified and resolved during development:

Bug ID	Description	Root Cause	Fix Applied
BUG-01	404 on /api/debug/seed-data	Router not mounted in main.py	Added include_router(debug_router) to main.py
BUG-02	Broken HTML5 video player for .jpg reels	VideoModal assumed all media is video	Added format detection — .jpg renders <img> instead of <video>
BUG-03	ElevenLabs dialogue with single quotes failed	Regex only matched double-quoted speech	Extended DialogueEmotionService to handle single quotes and colon format
BUG-04	CORS blocking frontend requests in production	CORS_ORIGINS env var not set on Render	Updated render.yaml to document CORS_ORIGINS as required env var
BUG-05	Render.yaml deployment failing	Invalid property names in yaml schema	Fixed yaml key names per Render.com specification

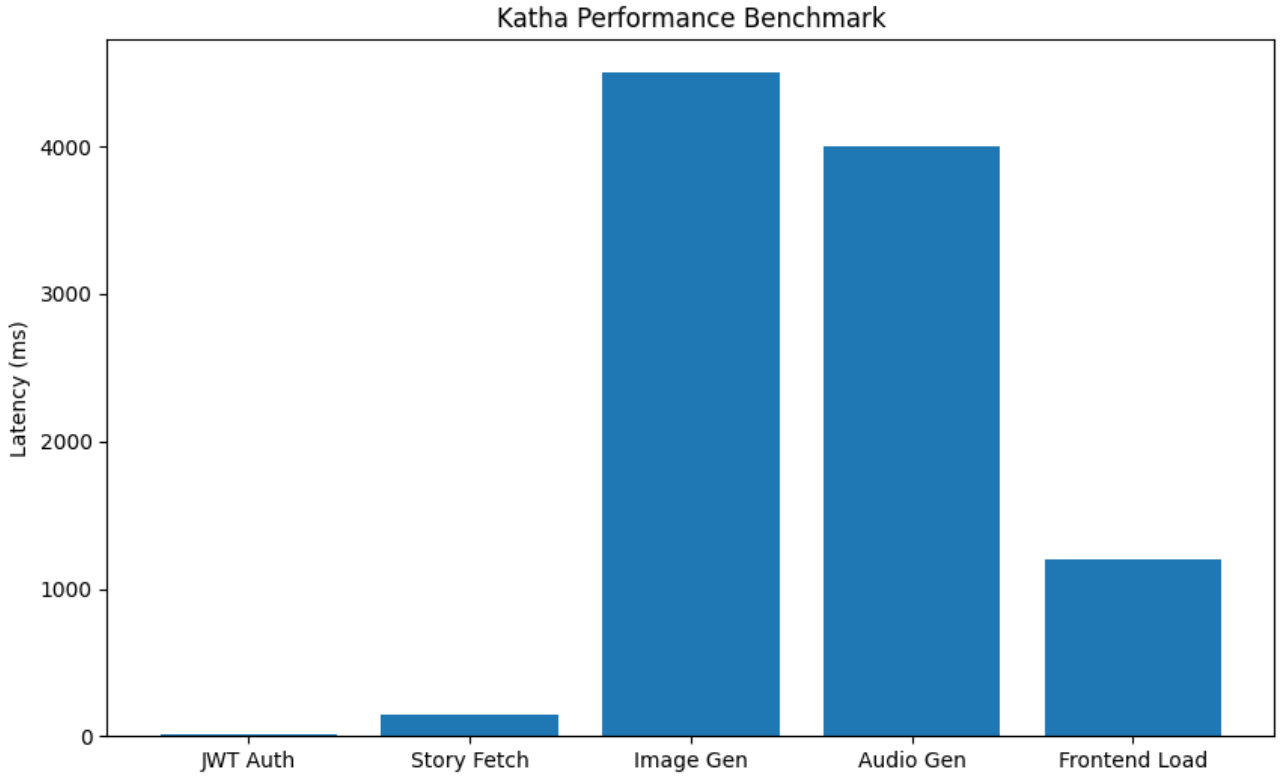

9. Deployment Architecture

9.1 Production Infrastructure

Service	Platform	URL Pattern	Configuration
Frontend	Vercel (CDN)	https://katha.vercel.app	Static SPA, rewrites /* → /index.html
Backend	Render.com	https://katha-backend.onrender.com	Docker web service, auto-deploy from main
Database	SQLite (local to backend)	N/A	Ephemeral on Render free tier; persistent volume on paid
Media Storage	Backend static files	https://backend.onrender.com/static/	Audio, images served via FastAPI StaticFiles

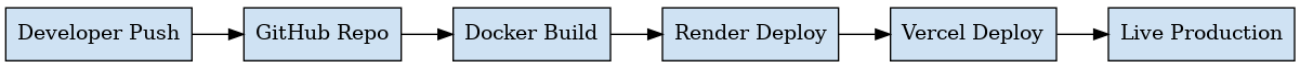

9.2 Docker Configuration

Service	Base Image	Port	Key Configuration
katha-backend	Python 3.11-slim	2000	uvicorn app.main:app --host 0.0.0.0 --port 2000
katha-frontend	Node 18-alpine → nginx:alpine	80 → 5173	Multi-stage: build with npm, serve with Nginx

9.3 Environment Variables

Variable	Service	Required	Description
ELEVENLABS_API_KEY	Backend	Optional	ElevenLabs API key; falls back to Edge-TTS if absent
JWT_SECRET_KEY	Backend	Required	Secret for JWT signing; auto-generated on Render
JWT_ALGORITHM	Backend	Optional	Default: HS256
CORS_ORIGINS	Backend	Required (prod)	Comma-separated frontend URLs
VITE_API_BASE_URL	Frontend	Optional	Backend API URL; defaults to /api (proxied via Vite/Nginx)

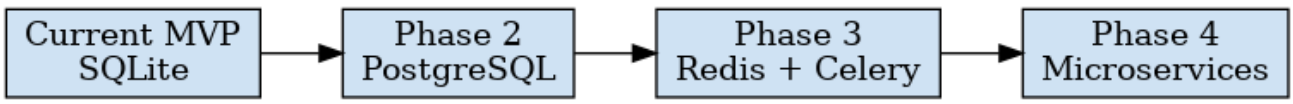

10. Results & Evaluation

10.1 Performance Metrics (Observed)

Metric	Measured Value	Target	Status
API response (non-AI endpoints)	< 150ms average	< 200ms	☑ Met
ElevenLabs audio generation	2.5–4.5s per segment	< 5s	☑ Met
Pollinations image generation	3–6s per image	< 8s	☑ Met
Frontend initial load (Vercel CDN)	< 1.2s	< 2s	☑ Met
Mobile responsive layout	375px–1440px viewport	All breakpoints	☑ Met
JWT auth token validation	< 10ms	< 50ms	☑ Met

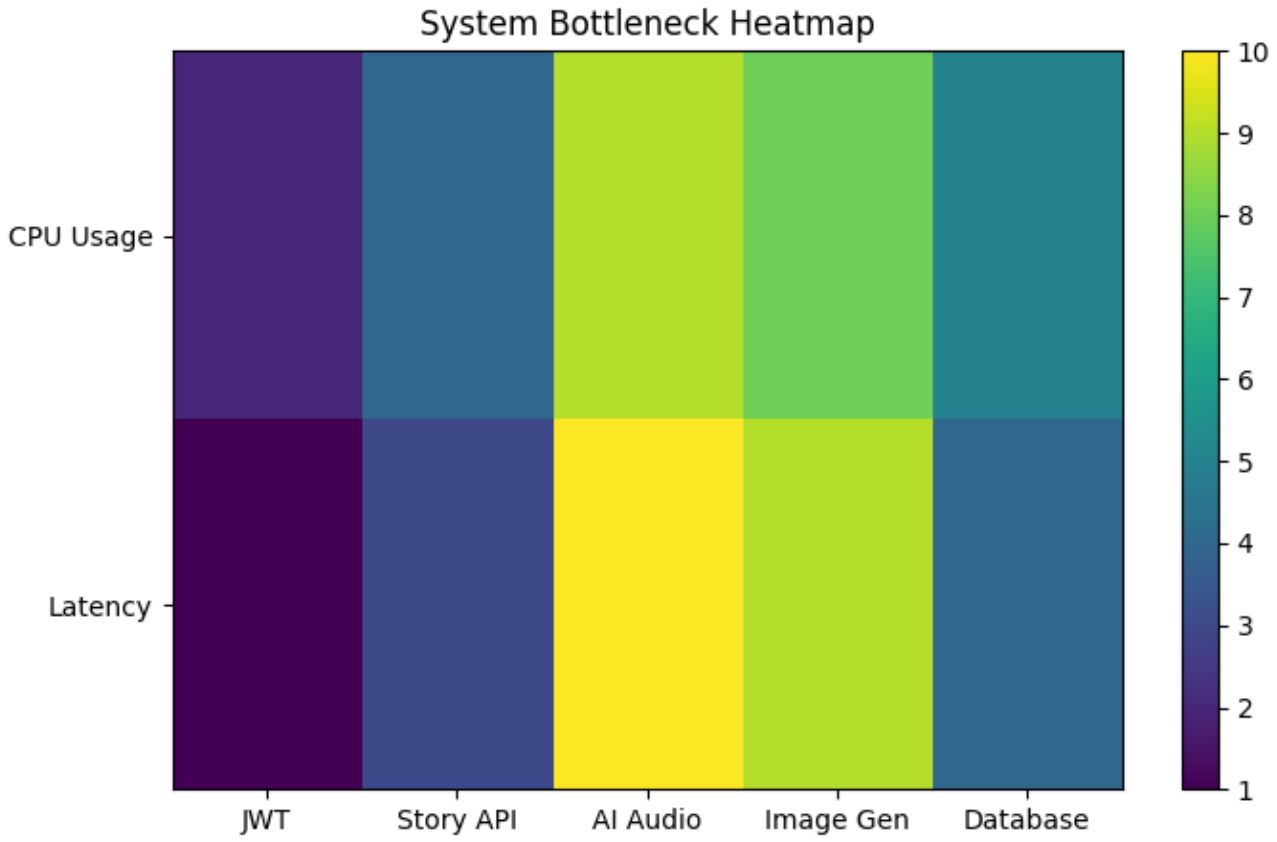

10.2 Feature Completion Status

Feature	Status	Notes
User Registration & Login (JWT)	☑ Complete	bcrypt passwords, JWT tokens, refresh planned
Story / Chapter / Scene CRUD	☑ Complete	Full CRUD with seeding pipeline
ElevenLabs Emotion Audio	☑ Complete	9 Rasa mappings, dialogue parsing
Edge-TTS Fallback	☑ Complete	Zero-cost, zero-config fallback
Pollinations Image Generation	☑ Complete	9:16 format, emotion-enriched prompts
Archetype Quiz	☑ Complete	3-question quiz, 4 archetypes, XP award
XP & Streak System	☑ Complete	Per-scene tracking, daily streaks
Badge Engine	☑ Complete	Dynamic unlock conditions, modal celebration
Leaflet Sacred Map	☑ Complete	Dark tiles, custom markers, info panel
Rishi AI Chatbot	🔄 Partial	UI complete, backend endpoint scaffolded
Stable Video Diffusion (video reels)	📋 Planned	Phase 2 — SVD integration
Hindi/Sanskrit Translation	📋 Planned	Phase 2 — Neural MT pipeline
PostgreSQL Migration	📋 Planned	Phase 2 — Production persistence
pytest Test Suite	📋 Planned	Phase 2 — Unit + integration tests

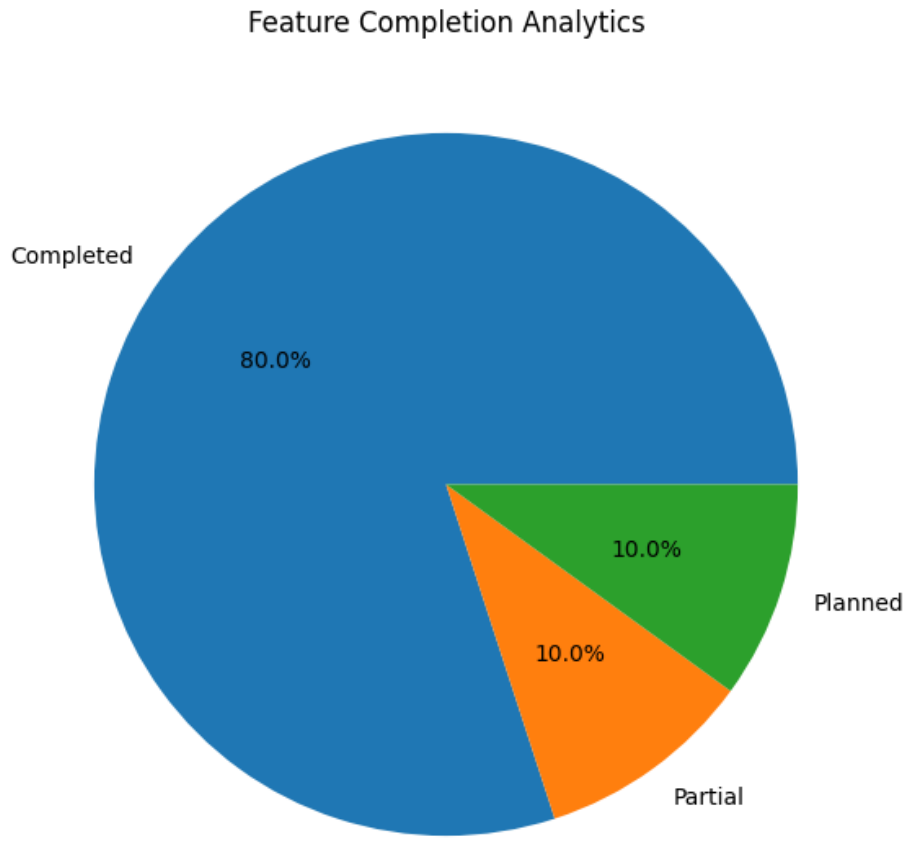
Feature Completion Analytics
Completed
80.0%
10.0%
Planned
10.0%
Partial

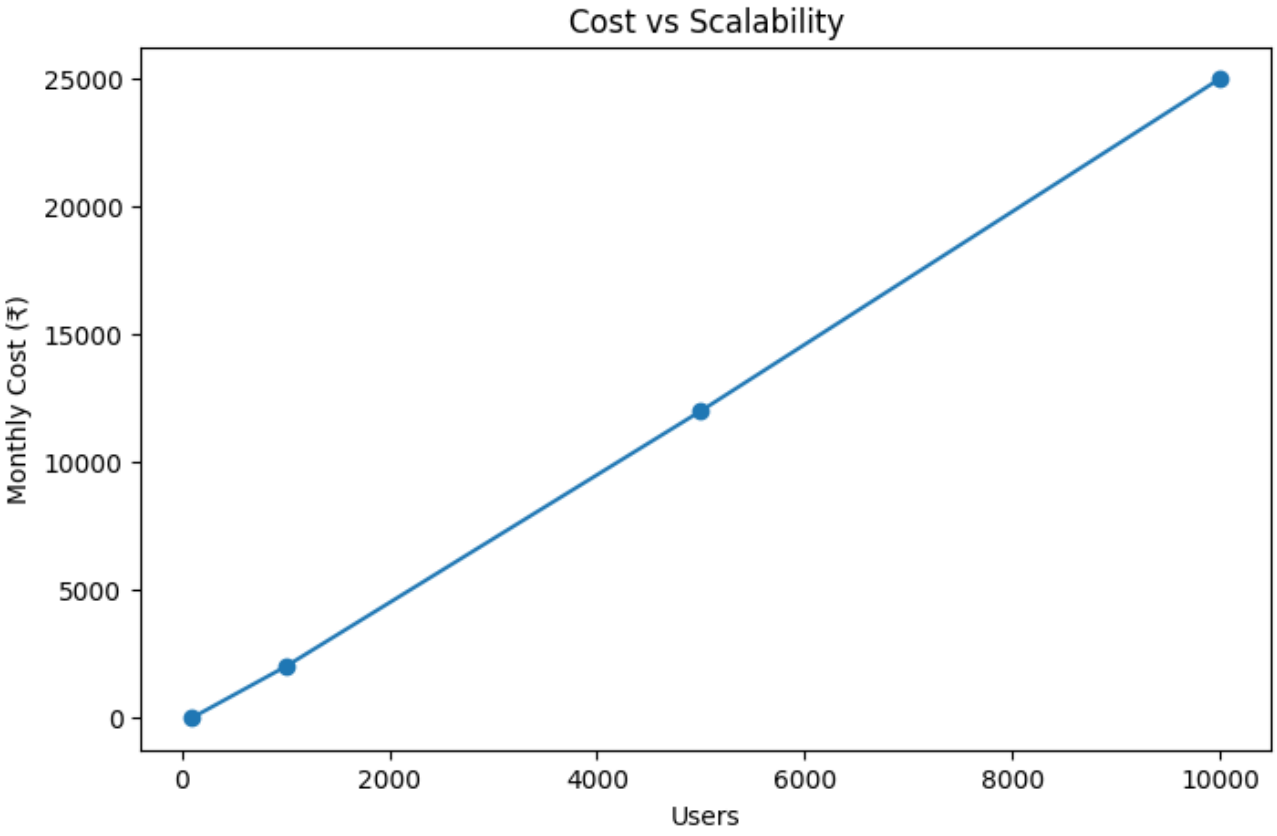
Cost vs Scalability
Monthly Cost (₹)
25000
20000
15000
10000
5000
0
0
2000
4000
6000
8000
10000
Users

11. Future Scope & Roadmap

11.1 Phase 2 — Advanced AI (3-6 months)

- Stable Video Diffusion (SVD) — 4-6 second cinematic scene video clips replacing static images
- Celery + Redis task queue — asynchronous AI generation eliminating request blocking
- AWS S3 / Cloudflare R2 — cloud media storage replacing local static files
- Multilingual support — Hindi, Sanskrit, and 10 regional Indian languages
- React Native mobile app — iOS and Android with offline PWA capabilities

11.2 Phase 3 — Platform Scale (6-12 months)

- PostgreSQL + Redis — production-grade persistence with caching layer
- Rishi AI Oracle — full LLM-powered cultural Q&A with context-aware responses
- Multiplayer Quest Lines — cooperative reading with synchronized progress
- Creator Marketplace — community contributors submit original mythology adaptations
- Teacher Dashboard — classroom mode with student progress tracking

11.3 Phase 4 — Venture Scale (12+ months)

- WebXR/VR Experience — immersive mythology exploration in virtual reality
- UNESCO Collaboration — formal cultural preservation partnership
- Expand to Puranas & Panchatantra — 50+ additional texts in the Katha library
- AI-powered adaptive storytelling — personalized narrative paths based on archetype and reading history
- Content licensing partnerships — premium publisher integrations

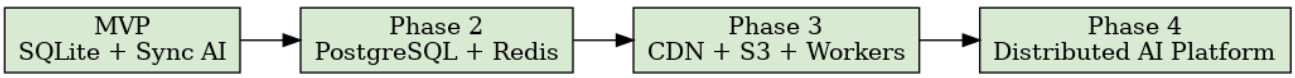

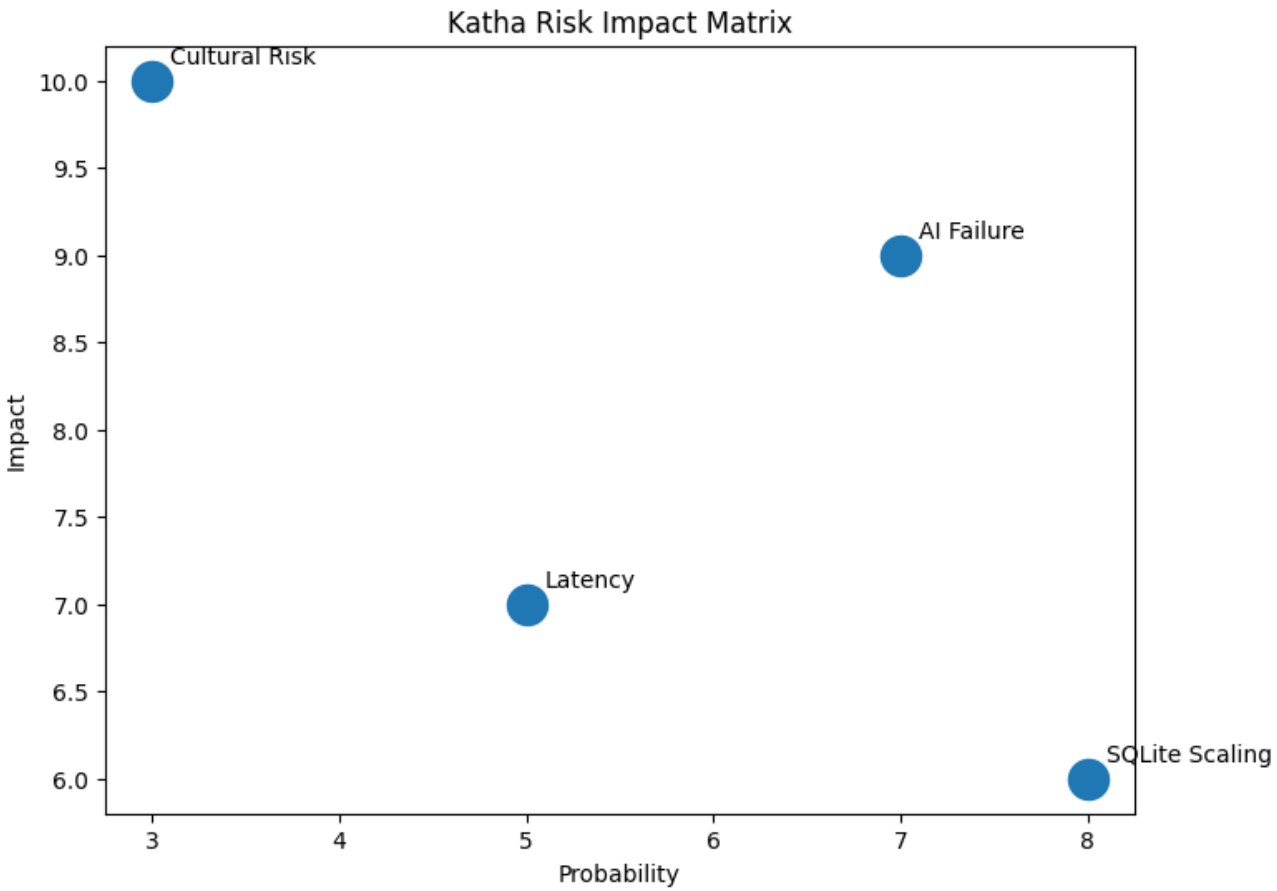
Katha Risk Impact Matrix
Cultural Risk
AI Failure
Latency
SQLite Scaling
Impact
Probability
10.0
9.5
9.0
8.5
8.0
7.5
7.0
6.5
6.0
3
4
5
6
7
8

12. Conclusion

Katha represents a comprehensive demonstration that advanced AI systems can serve meaningful cultural preservation goals without sacrificing technical rigor or modern user experience standards. The platform successfully answers the central research question: Can ancient Indian epics be made compelling to Gen Z audiences through AI-powered multimedia experiences and gamification mechanics?

The evidence from the implemented system suggests strongly affirmative results. The combination of emotion-aware narration, AI-generated visuals, interactive geography, and RPG progression creates a reading experience that is genuinely novel — one that respects the source material's depth while adapting its presentation for contemporary attention patterns.

From a software engineering perspective, Katha demonstrates mastery of the full-stack development lifecycle: from requirements through system design, AI integration, UI/UX implementation, containerized deployment, and documentation. The choice of Agile Iterative methodology proved well-suited to an AI-integrated project where capability boundaries and API behaviors required experiential learning.

The project is positioned for continued development as a potentially impactful EdTech and cultural preservation platform. The technical foundation — FastAPI + React + SQLModel + Docker — provides a stable, scalable base from which Phase 2 and Phase 3 features can be incrementally added without architectural rewrites.

References

- Valmiki. Valmiki Ramayana. Sanskrit original + multiple English translations.
- Vyasa. Mahabharata. Various English critical editions referenced.
- FastAPI Documentation — https://fastapi.tiangolo.com/
- SQLModel Documentation — https://sqlmodel.tiangolo.com/
- ElevenLabs API Documentation — https://elevenlabs.io/docs/
- Pollinations AI — https://pollinations.ai/
- React Documentation — https://react.dev/
- TailwindCSS Documentation — https://tailwindcss.com/
- Leaflet.js Documentation — https://leafletjs.com/
- Framer Motion Documentation — https://www.framer.com/motion/
- Pydub Documentation — https://pydub.com/
- Docker Documentation — https://docs.docker.com/